

Pytorch basics

Adapted from Day 3 Lab 1

by yours truly, farah

Table of contents

- Intro to tensors
- Data representation in Deep Learning
- Tensor dimension operations
- Tensor Data types and transfer to GPUs
- Pytorch workflow
 - Dataset class
 - Model class
 - Training loop
 - Validation loop

why Pytorch?

- Dynamic, changes occur during run-time
- Flexible for fast prototyping
- Often used in research experimentation

- Tensorflow is for old people
- The only reason I would ever use tensorflow:





my GOAT Andrew

Pro tip: Use Andrew Ng's course on Neural Networks and Deep Learning offered to you on Coursera

PyTorch is an open-source deep learning framework that allows us to build and train neural networks using:

tensors and automatic differentiation.



what even are tensors?

- Tensors are a way of storing the input data
- **In most cases**, they are multi-dimensional matrices
- They contain elements of a SINGLE DATA TYPE, just like:
 - `torch.int32`
 - `torch.float64`
- Why tensors?
 - Generalize data of all dimensions (0D, 1D, 2D... 6D)
 - They run on both CPUs and GPUs

what even are tensors?

Scaler	Vector	Matrix	Tensor
1	$\begin{bmatrix} 7 \\ 8 \end{bmatrix}$	$\begin{bmatrix} 7 & 2 \\ 8 & 4 \end{bmatrix}$	$\begin{bmatrix} \begin{bmatrix} 7 & 2 \\ 8 & 4 \end{bmatrix} & \begin{bmatrix} 2 & 7 \\ 4 & 8 \end{bmatrix} \end{bmatrix}$

question: how many tensors are present in the image above?

answer: 4, they're all tensors, just different dimensions

creating tensors

```
import torch

# Create tensors
x = torch.tensor([1.0, 2.0, 3.0])
y = torch.randn(3)
print("x:", x)
print("y:", y)

x: tensor([1., 2., 3.])
y: tensor([0.3548, 0.3962, 0.4785])
```

tensor shapes

```
print("Shape of x:", x.shape)
```

```
Shape of x: torch.Size([3])
```

tensor operations

given:

```
a = torch.tensor([1.0, 2.0, 3.0])  
b = torch.tensor([4.0, 5.0, 6.0])
```

elementwise operations:

```
print("Addition:", a + b)
```

Addition: tensor([5., 7., 9.])

```
print("Multiplication:", a * b)
```

Multiplication: tensor([4., 10., 18.])

tensor operations

matrix multiplication:

```
A = torch.randn(2, 3)  
B = torch.randn(3, 2)  
C = A @ B
```

A: tensor([[0.0721, -0.9560, -1.2007],
[-0.1852, -0.4110, -0.9032]])

×

B: tensor([[0.0721, -0.9560, -1.2007],
[-0.1852, -0.4110, -0.9032]])

C: tensor([[-1.0382, -0.5640],
[-0.0658, 0.5422]])

Matrix multiplication result shape: torch.Size([2, 2])

data representation

The way data is represented in tensors depends on whether it is:

structured


Tabular Data

columns = attributes for those observations

Rows = observations

Player	Minutes	Points	Rebounds	Assists
A	41	20	6	5
B	30	29	7	6
C	22	7	7	2
D	26	3	3	9
E	20	19	8	0
F	9	6	14	14
G	14	22	8	3
I	22	36	0	9
J	34	8	1	3

unstructured



structured data

- represented as a 2D tensor
- Shape: (batch_size, number_of_features)
- Commonly used for tasks like regression and classification

question: what is the tensor shape of the table on the right?

answer: (9, 5)

Tabular Data

columns = attributes for those observations

Player	Minutes	Points	Rebounds	Assists
A	41	20	6	5
B	30	29	7	6
C	22	7	7	2
D	26	3	3	9
E	20	19	8	0
F	9	6	14	14
G	14	22	8	3
I	22	36	0	9
J	34	8	1	3


Rows = observations

unstructured data

- represented as a 4D tensor
- Shape: (batch_size, channels, height, width)
- Channels represent color information:
 - Grayscale → 1 channel
 - RGB (Red, Green, Blue) → 3 channels

question: assume the image on the right is composed of 1024x1024 pixels, and we have 120 images in total, what is the tensor shape?

answer: (120, 3, 1024, 1024)



tensor shape manipulation

- Manipulating tensor shapes is essential in deep learning
- I want you to know these functions better than you know your own family:
 - `.flatten()`
 - `.squeeze()`
 - `.unsqueeze()`
 - `.view()`

.flatten()

- Flattens input by reshaping it into a 1D tensor
- Converts **any shape** to (batch_size, features)



.flatten()

```
import torch

x = torch.randn(32, 3, 28, 28)
x_flat = x.flatten(start_dim=1)
print("Flatten:", x_flat.shape)
```

Flatten: torch.Size([32, 2352])

***note:** you can use both x.flatten() or torch.flatten(x)

.flatten()

torch.flatten

`torch.flatten(input, start_dim=0, end_dim=-1) → Tensor`

Flattens `input` by reshaping it into a one-dimensional tensor. If `start_dim` or `end_dim` are passed, only dimensions starting with `start_dim` and ending with `end_dim` are flattened. The order of elements in `input` is unchanged.

Unlike NumPy's flatten, which always copies input's data, this function may return the original object, a view, or copy. If no dimensions are flattened, then the original object `input` is returned. Otherwise, if input can be viewed as the flattened shape, then that view is returned. Finally, only if the input cannot be viewed as the flattened shape is input's data copied. See [torch.Tensor.view\(\)](#) for details on when a view will be returned.

 Note

Flattening a zero-dimensional tensor will return a one-dimensional view.

Parameters:

- `input` ([Tensor](#)) – the input tensor.
- `start_dim` ([int](#)) – the first dim to flatten
- `end_dim` ([int](#)) – the last dim to flatten

.flatten()

question: given a tensor shape of (32, 3, 28, 28), what is the output shape if we flatten the tensor with the start_dim = 0?

answer: 75264

question: with the same tensor, what if the start_dim = 2?

answer: (32, 3, 784)

.flatten() common enemies (1)

RuntimeError: mat1 and mat2 shapes cannot be multiplied

why does this happen?

- Fully connected layers expect 2D tensors: (batch_size, features)
- But if you forget to flatten, you might pass something like:
(batch_size, channels, height, width)

Fix: `x.flatten(start_dim=1)`

.flatten() common enemies (2)

RuntimeError: Expected input size (N, C) but got (N, d1, d2)

why does this happen?

Predictions are multi-dimensional (1D or 2D targets)

Fix: `x.flatten(start_dim=1)`

.squeeze()

- Removes dimensions with size 1
- For example, if `input` is of shape: $(A \times 1 \times B \times C \times 1 \times D)$ then `input.squeeze()` will be of shape: $(A \times B \times C \times D)$



.squeeze()

```
import torch

x = torch.randn(1, 3, 28, 28)
x_sq = x.squeeze()
print("Squeeze:", x_sq.shape)
```

Squeeze: torch.Size([3, 28, 28])

***note:** you can use both x.squeeze() or torch.squeeze(x)

.squeeze()

torch.squeeze

`torch.squeeze(input: Tensor, dim: Optional[Union[int, List[int]]]) → Tensor`

Returns a tensor with all specified dimensions of `input` of size 1 removed.

For example, if `input` is of shape: $(A \times 1 \times B \times C \times 1 \times D)$ then the `input.squeeze()` will be of shape: $(A \times B \times C \times D)$.

When `dim` is given, a squeeze operation is done only in the given dimension(s). If `input` is of shape: $(A \times 1 \times B)$, `squeeze(input, 0)` leaves the tensor unchanged, but `squeeze(input, 1)` will squeeze the tensor to the shape $(A \times B)$.

Note

The returned tensor shares the storage with the input tensor, so changing the contents of one will change the contents of the other.

Warning

If the tensor has a batch dimension of size 1, then `squeeze(input)` will also remove the batch dimension, which can lead to unexpected errors. Consider specifying only the dims you wish to be squeezed.

.squeeze()

question: given a tensor x of shape $(1, 3, 28, 28)$, what is the output shape when we apply $x.\text{squeeze}()$?

answer: $(3, 28, 28)$

question: given a tensor x of shape $(2, 1, 2, 1, 2, 1)$, what is the output shape when we apply $x.\text{squeeze}()$?

answer: $(2, 2, 2)$

.squeeze() common enemy

RuntimeError: Expected 2D input, got 1D tensor

This usually happens during:

- Validation
- Inference
- Debugging with a single sample

Fix: `x.squeeze(dim=1)`

.unsqueeze()

- Adds a dimension with size 1 at a specified position.
- The position is controlled through the `dim` index



.unsqueeze()

```
import torch

x = torch.randn(3, 28, 28)
x_unsq = x.unsqueeze(0)
print("Unsqueeze:", x_unsq.shape)
```

Unsqueeze: torch.Size([1, 3, 28, 28])

***note:** you can use both x.unsqueeze() or torch.unsqueeze(x)

.unsqueeze()

torch.unsqueeze

`torch.unsqueeze(input, dim)` → [Tensor](#)

Returns a new tensor with a dimension of size one inserted at the specified position.

The returned tensor shares the same underlying data with this tensor.

A `dim` value within the range `[-input.dim() - 1, input.dim() + 1]` can be used. Negative `dim` will correspond to [unsqueeze\(\)](#) applied at `dim = dim + input.dim() + 1`.

Parameters:

- `input` ([Tensor](#)) – the input tensor.
- `dim` ([int](#)) – the index at which to insert the singleton dimension

.unsqueeze()

question: given the tensor `x = torch.tensor([1, 2, 3, 4])`, what is the output tensor when we perform `x.unsqueeze(dim=0)` ?

answer: `tensor([[1, 2, 3, 4]])`

question: what is the output tensor if we instead apply `x.unsqueeze(dim=1)` ?

answer: `tensor([[1], [2], [3], [4]])`

.unsqueeze() common enemies

RuntimeError: Expected input batch_size (N) to match target batch_size (M)


or

RuntimeError: mat1 and mat2 shapes cannot be multiplied

Fix: `x.unsqueeze(dim=1)` and check tensor shapes

.view()

- Reshapes a tensor freely while maintaining the same number of elements.
- Input parameter ***shape** is to control the desired size



.view()

```
import torch

x = torch.randn(32, 28, 28, 3)
x_view = x.view(32, -1)
print("View:", x_view.shape)
```

View: torch.Size([32, 2352])

*Important note: In view, -1 infers the dimension automatically so that the total number of elements stays the same.

Think -1 = whatever size makes the math work

.view()

torch.Tensor.view

`Tensor.view(*shape) → Tensor`

Returns a new tensor with the same data as the `self` tensor but of a different `shape`.

The returned tensor shares the same data and must have the same number of elements, but may have a different size. For a tensor to be viewed, the new view size must be compatible with its original size and stride, i.e., each new view dimension must either be a subspace of an original dimension, or only span across original dimensions $d, d + 1, \dots, d + k$ that satisfy the following contiguity-like condition that $\forall i = d, \dots, d + k - 1$,

$$\text{stride}[i] = \text{stride}[i + 1] \times \text{size}[i + 1]$$

Otherwise, it will not be possible to view `self` tensor as `shape` without copying it (e.g., via `contiguous()`).

When it is unclear whether a `view()` can be performed, it is advisable to use `reshape()`, which returns a view if the shapes are compatible, and copies (equivalent to calling `contiguous()`) otherwise.

Parameters:

`shape (torch.Size or int...)` – the desired size

.view()

question: given the tensor `x = torch.randn(4, 4)`, what is the output shape when we perform `x.view(16)`?

answer: 16

question: what is the output shape if we instead performed `x.view(-1, 8)`?

answer: (2, 8)

Changing data type

Depending on the loss function used, prediction and target tensors need to be cast to a specific dtype

Examples:

- nn.BCEWithLogitsLoss requires:
 - Predictions: float
 - Targets: float
- If targets are int, we will get:

RuntimeError: result type Float can't be cast to Long

Changing data type

Depending on the loss function used, prediction and target tensors need to be cast to a specific dtype

question: What are the appropriate data types for **logits** and **targets** when computing the loss using nn.CrossEntropyLoss?

*hint: Go to nn.CrossEntropyLoss documentation, and look at the examples section

answer: logits: float, targets: long

Changing data type

- Use `.to(dtype)` to convert a tensor's data type.

```
import torch

# Create a float32 tensor
x = torch.tensor([1.2, 2.3, 3.4], dtype=torch.float32)
print(x.dtype) # Output: torch.float32

# Convert to float16
x_half = x.to(torch.float16)
print(x_half.dtype) # Output: torch.float16

torch.float32
torch.float16
```

Move Tensors to GPU (if available)

- GPUs are faster and more efficient in most cases when training or inferencing deep learning models
- Use `.to(device)` to move a tensor to GPU for faster computation.

```
# Automatically select CPU or GPU
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Create a tensor and move it to GPU
x_gpu = x.to(device)
print(x_gpu.device) # Output: cuda:0 (if GPU is available) or cpu

cpu
```

PyTorch Workflow Organization

consists of 4 main components:

1

Dataset
Class

2

Model Class

3

Training
Loop

4

Validation
Loop

Dataset class

- A Dataset Class is responsible for transforming raw data into samples that are ready to be used by a model.
- Each sample returned consists of:
 - An input (features or image)
 - Its corresponding label

Dataset class

Typical workflow:

1. Load data
2. Train/Test split
3. Standardize
4. Transform to tensors
5. Wrap in TensorDataset
6. Use DataLoader

Dataset class

1.

```
# load data
data = load_breast_cancer()
X = data.data
y = data.target
```

2.

```
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

3.

```
# Standardize the data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

4.

```
# transform to tensors
X_train = torch.tensor(X_train, dtype=torch.float32)
X_test = torch.tensor(X_test, dtype=torch.float32)
y_train = torch.tensor(y_train, dtype=torch.float32)
y_test = torch.tensor(y_test, dtype=torch.float32)
```

Dataset class

5.

```
# TensorDataset pairs input features (X) with their corresponding labels (y)
# Each item in the dataset is returned as (X[i], y[i])
train_dataset = TensorDataset(X_train, y_train)
test_dataset = TensorDataset(X_test, y_test)
```

6.

```
# DataLoader for training data
train_loader = DataLoader(
    train_dataset,
    batch_size=32,
    shuffle=True,
    num_workers=2
)
# DataLoader for test/validation data
test_loader = DataLoader(
    test_dataset,
    batch_size=32,
    shuffle=False,
    num_workers=2
)
```


TensorDataset

- A wrapper that lets you treat one or more tensors as a dataset, so they can be fed cleanly into a `DataLoader` for training.
- we can use `TensorDataset` to pair input features with their labels.

DataLoader

Problem: We need batches, not single samples

Solution: We use DataLoader to handle batching automatically.



DataLoader

- A `DataLoader` is a PyTorch utility that takes a `Dataset` and does:
 - **Batching**: Groups multiple samples together for faster processing.
 - **Shuffling**: Randomizes data order to improve training.
 - **Multi-threading**: Loads data efficiently in parallel.

Argument	Description
<code>dataset</code>	The dataset object (e.g., <code>train_dataset</code>)
<code>batch_size</code>	Number of samples per batch (e.g., <code>32</code>)
<code>shuffle</code>	Whether to randomly shuffle data each epoch (<code>True</code> = better for training)
<code>num_workers</code>	Number of parallel CPU workers to load data faster
<code>collate_fn</code>	A function to customize how data is stacked (useful when data has variable sizes)

Model Class

In PyTorch, `nn.Linear(in_features, out_features)` creates a fully connected (**dense**) layer that applies a linear transformation:

$$y = xW^T + b$$

`in_features`: number of input features

`out_features`: number of neurons (outputs) in the layer

Model Class

```
import torch
import torch.nn as nn

# Create a linear layer: 3 input features -> 1 output
linear = nn.Linear(in_features=3, out_features=1)

# Example input: batch of 4 samples, each with 3 features
x = torch.randn(4, 3)

# Forward through the layer
y = linear(x)

print("Input shape:", x.shape)    # torch.Size([4, 3])
print("Output shape:", y.shape)   # torch.Size([4, 1])
```

Input shape: torch.Size([4, 3])
Output shape: torch.Size([4, 1])


Model Class

Inside the `__init__` method, we define what the neural network looks like. This includes:

- Number of layers (`nn.Linear`)
- Hidden layer sizes
- Activation functions (ex: ReLU)
- Output activation function (ex: Sigmoid, Softmax)

Model Class

The `forward()` method defines how the input data flows through the model to produce the final output.



Model Class

1-Layer NN

```
class NN1Layer(nn.Module):

    def __init__(self, input_dim):

        super(NN1Layer, self).__init__()
        # input_dim = num of features, Output for binary classification is 1
        self.layer_1 = nn.Linear(input_dim, 1)

        # output activation function
        self.sigmoid = nn.Sigmoid()

    # forward pass
    def forward(self, x):
        z = self.layer_1(x)
        a = self.sigmoid(z)
        return a
```

2-Layer NN

```
class NN2Layer(nn.Module):

    def __init__(self, input_dim, hidden_dim):

        super(NN2Layer, self).__init__()
        # input_dim = num of features, hidden_dim = num of neurons
        self.layer1 = nn.Linear(input_dim, hidden_dim)
        # hidden_dim = num of neurons, Output for binary classification is 1
        self.layer2 = nn.Linear(hidden_dim, 1)

        # non-linearity activation function
        self.relu = nn.ReLU()
        # output activation function
        self.sigmoid = nn.Sigmoid()

    # forward pass
    def forward(self, x):
        z1 = self.layer1(x)
        a1 = self.relu(z1)
        z2 = self.layer2(a1)
        a2 = self.sigmoid(z2)
        return a2
```

Training Loop

The training loop is responsible for updating the model's weights so that it learns to minimize the loss function.

Parameters

- model – The neural network to be trained.
- optimizer – Updates model parameters (e.g., SGD, Adam).
- criterion – Loss function, depends on the task.
- train_loader – PyTorch DataLoader that provides batches of training data.
- device – Device used for computation (cpu or cuda).

Training Loop

```
def train_one_epoch(model, optimizer, criterion, train_loader, device):
    # Set the model to training mode
    model.train()

    running_loss = 0.0

    for X_batch, y_batch in train_loader:
        # Move batch to the selected device
        X_batch = X_batch.to(device)
        y_batch = y_batch.view(-1, 1).to(device)

        # Forward pass
        outputs = model(X_batch)
        loss = criterion(outputs, y_batch)

        # Backward pass & optimization
        optimizer.zero_grad()    # Clear previous gradients
        loss.backward()           # Compute gradients
        optimizer.step()          # Update model parameters

        running_loss += loss.item()

    # Average loss over all batches
    avg_loss = running_loss / len(train_loader)

    return avg_loss
```

Validation Loop

The validation loop evaluates the model's performance on unseen data without updating the weights.

It is used to measure how well the model generalizes.

Parameters

- `model` – The trained neural network to be evaluated.
- `criterion` – Loss function used for evaluation
- `test_loader` – PyTorch `DataLoader` that provides batches of validation/test data.
- `device` – Device used for computation (`cpu` or `cuda`).

Validation Loop

```
def validate(model, criterion, test_loader, device):
    # Set the model to evaluation mode
    model.eval()

    running_loss = 0.0
    correct = 0
    total = 0

    with torch.no_grad(): # Disable gradient computation
        for X_batch, y_batch in test_loader:
            # Move batch to the selected device
            X_batch = X_batch.to(device)
            y_batch = y_batch.view(-1, 1).to(device)

            # Forward pass
            outputs = model(X_batch)
            loss = criterion(outputs, y_batch)
            running_loss += loss.item()

            # Binary predictions
            predicted = (outputs > 0.5).float()

            # Accuracy calculation
            correct += (predicted == y_batch).sum().item()
            total += y_batch.size(0)

    avg_loss = running_loss / len(test_loader)
    accuracy = correct / total

    return avg_loss, accuracy
```

Note: Gradients are disabled during validation using `torch.no_grad()` to improve efficiency and prevent weight updates.